

HW2 by Trisha Mandal

```
In [1]: import pandas as pd
import numpy as np
import nltk
nltk.download('wordnet')
import re
from bs4 import BeautifulSoup
from textblob import TextBlob
from sklearn.model_selection import GridSearchCV
import warnings
from sklearn.model_selection import train_test_split
warnings.filterwarnings('ignore')
```

```
[nltk_data] Downloading package wordnet to
[nltk_data] /Users/trishamandal/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

```
In [2]: pip install torch
```

```
Requirement already satisfied: torch in /Users/trishamandal/opt/anaconda3/lib/
python3.9/site-packages (1.12.1)
Requirement already satisfied: typing-extensions in /Users/trishamandal/opt/an
aconda3/lib/python3.9/site-packages (from torch) (4.1.1)
Note: you may need to restart the kernel to use updated packages.
```

1. Dataset Generation

We will use the Amazon reviews dataset used in HW1. Load the dataset and build a balanced dataset of 100K reviews along with their ratings to create labels through random selection. You can store your dataset after generation and reuse it to reduce the computational load. For your experiments consider a 80%/20% training/testing split.

```
In [3]: # reading the data
df = pd.read_table('amazon_reviews_us_Jewelry_v1_00.tsv', error_bad_lines=False)
```

```
In [4]: # selecting only reviews and ratings
data = pd.concat([df['star_rating'], df['review_body']], axis=1)
```

```
In [5]: #converting all reviews into string
data['review_body'] = [str(i) for i in data['review_body']]
```

```
In [6]: # data cleaning on sampled data
# step 1: changing all words to lower case by using str.lower()
# step 2: performing contractions on the reviews by using contractions library
# step 3: Removing HTML by using BeautifulSoup library
# step 4: Removing URLs by using regex
# step 5: Removing non-alphanumeric characters by using regex
# step 6: Stripping extra space
# step 7: Replacing double spaces with single spaces
import contractions
```

```

data.dropna()
data['review_body'] = data['review_body'].str.lower()
# data.drop(data[data['review_body'].str.split().str.len() < 10].index, inplace=True)
data['review_body'] = data['review_body'].apply(lambda x: contractions.fix(x))
data['review_body'] = [BeautifulSoup(text).get_text() for text in data['review_body']]
data['review_body'] = data['review_body'].apply(lambda text: re.sub(r'www.\S+', ''))
data['review_body'] = data['review_body'].apply(lambda text: re.sub(r'https?://\S+', ''))
data['review_body'] = data['review_body'].apply(lambda x: re.sub('\W+', ' ', x))
data['review_body'] = data['review_body'].apply(lambda x: re.sub(r'\d+', ' ', x))
data['review_body'] = data['review_body'].apply(lambda x: x.strip())
data['review_body'] = data['review_body'].apply(lambda x: x.replace(" ", ""))

```

```

In [8]: #seperating reviews from each rating
star1 = data[data['star_rating'] == 1]
star2 = data[data['star_rating'] == 2]
star3 = data[data['star_rating'] == 3]
star4 = data[data['star_rating'] == 4]
star5 = data[data['star_rating'] == 5]
star3

```

```

Out[8]:

```

	star_rating	review_body
17	3	not what i expected took too long to ship but ...
35	3	it states that the item is new but you cannot ...
40	3	the bracelet did not fit properly i had to act...
83	3	i have never been able to use these on my ring...
88	3	arrived broken one if the stones were out look...
...
1766904	3	the braclet and ring set of infinite potential...
1766930	3	well i got this necklace and come on its ctw...
1766952	3	less is best i may buy the jewelry but i am no...
1766955	3	the earrings are very pretty but they were too...
1766962	3	just wanted to let buyers know that these earr...

153665 rows × 2 columns

```

In [10]: #calculated review length for each document and putting it in length column in
numofwords1 = star1["review_body"].apply(lambda x: len(str(x).split(' ')))
star1["length"] = pd.DataFrame(numofwords1)
numofwords2 = star2["review_body"].apply(lambda x: len(str(x).split(' ')))
star2["length"] = pd.DataFrame(numofwords2)
numofwords3 = star3["review_body"].apply(lambda x: len(str(x).split(' ')))
star3["length"] = pd.DataFrame(numofwords3)
numofwords4 = star4["review_body"].apply(lambda x: len(str(x).split(' ')))
star4["length"] = pd.DataFrame(numofwords4)
numofwords5 = star5["review_body"].apply(lambda x: len(str(x).split(' ')))
star5["length"] = pd.DataFrame(numofwords5)
star5

```

Out [10]:

	star_rating	review_body	length
0	5	so beautiful even though clearly not high end ...	32
1	5	great product i got this set for my mother as ...	72
2	5	exactly as pictured and my daughter s friend l...	32
3	5	love it fits great super comfortable and neat ...	18
4	5	got this as a mother s day gift for my mom and...	22
...
1766982	5	i love these earrings my boyfriend got me a pai...	39
1766983	5	not too much money but would make a good impre...	10
1766985	5	i was so impressed with this piece i am a jewe...	84
1766990	5	the kt gold earrings look remarkable would def...	14
1766991	5	it will be a gift to my special friend we know...	46

1041018 rows × 3 columns

In [11]:

```
#sorting datframes in descending order by review length
star1 = star1.sort_values(by='length', ascending=False)
star2 = star2.sort_values(by='length', ascending=False)
star3 = star3.sort_values(by='length', ascending=False)
star4 = star4.sort_values(by='length', ascending=False)
star5 = star5.sort_values(by='length', ascending=False)
star1
```

Out [11]:

	star_rating	review_body	length
1648184	1.0	the entire review history is below but my most...	1489
1431879	1	the post broke off the face while trying to ge...	1139
1216795	1	i am writing a review for the three rings i pu...	1075
1757958	1	are of the sales profits of these red strings...	1035
1687894	1.0	i bought this bracelet as a christmas gift for...	952
...
740425	1	terrible	1
238902	1	garbage	1
743657	1	awful	1
278586	1	huge	1
884542	1	junk	1

150461 rows × 3 columns

In [12]:

```
#taking top 20K reviews from each dataframe
star1 = star1.head(20000)
star2 = star2.head(20000)
star3 = star3.head(20000)
```

```
star4 = star4.head(20000)
star5 = star5.head(20000)
star1
```

Out[12]:

	star_rating	review_body	length
1648184	1.0	the entire review history is below but my most...	1489
1431879	1	the post broke off the face while trying to ge...	1139
1216795	1	i am writing a review for the three rings i pu...	1075
1757958	1	are of the sales profits of these red strings...	1035
1687894	1.0	i bought this bracelet as a christmas gift for...	952
...
964376	1	i recieved this product a day earlier that is ...	65
1222447	1	i purchased this necklace to wear with an outf...	65
827380	1	i bought this because it looked like a great c...	65
1247483	1	i had bought this product based on the picture...	65
1680368	1.0	the image of the pendant is quite beautiful bu...	65

20000 rows × 3 columns

In [13]: *#combining all rating dataframes*

```
combined2 = pd.concat([star1, star2, star3, star4, star5])
combined2
```

Out[13]:

	star_rating	review_body	length
1648184	1.0	the entire review history is below but my most...	1489
1431879	1	the post broke off the face while trying to ge...	1139
1216795	1	i am writing a review for the three rings i pu...	1075
1757958	1	are of the sales profits of these red strings...	1035
1687894	1.0	i bought this bracelet as a christmas gift for...	952
...
1596083	5.0	i bought this ring as my wedding band thinking...	131
1715843	5	this ring is my new engagement wedding ring an...	131
1318493	5	this ring is gorgeous my diamond fell out of m...	131
1593443	5.0	i adore these earrings i also ordered the coor...	131
950317	5	this like the blue hearts opal ring i am cond...	131

100000 rows × 3 columns

In [109...]

```
combined2['review_body'] = [str(i) for i in combined2['review_body']]
combined2['star_rating'] = [int(i) for i in combined2['star_rating']]
```

In [14]: *# separating all reviews using ratings*

```

rating1 = data[data['star_rating'] == 1]
rating2 = data[data['star_rating'] == 2]
rating3 = data[data['star_rating'] == 3]
rating4 = data[data['star_rating'] == 4]
# done random sampling for rating 5 since number of reviews are too large
rating5 = data[data['star_rating'] == 5].sample(n=20000, random_state=

```

```

In [15]: # convertings reviews and ratings columns to lists for all rating categories
re1 = rating1['review_body'].values.tolist()
re2 = rating2['review_body'].values.tolist()
re3 = rating3['review_body'].values.tolist()
re4 = rating4['review_body'].values.tolist()

```

```

In [16]: rev1 = np.array(re1)
rev2 = np.array(re2)
rev3 = np.array(re3)
rev4 = np.array(re4)

```

```

In [17]: # applying tfidf on all ratings
from sklearn.feature_extraction.text import TfidfVectorizer
v1 = TfidfVectorizer()
v2 = TfidfVectorizer()
v3 = TfidfVectorizer()
v4 = TfidfVectorizer()
vec1 = v1.fit_transform(re1)
vec2 = v2.fit_transform(re2)
vec3 = v3.fit_transform(re3)
vec4 = v4.fit_transform(re4)

```

```

In [18]: #Converting to numpy
num1 = list(np.squeeze(np.asarray(np.sum(vec1, axis = 1).astype(np.float32))))
num2 = list(np.squeeze(np.asarray(np.sum(vec2, axis = 1).astype(np.float32))))
num3 = list(np.squeeze(np.asarray(np.sum(vec3, axis = 1).astype(np.float32))))
num4 = list(np.squeeze(np.asarray(np.sum(vec4, axis = 1).astype(np.float32))))

num1 = np.array(num1)
num2 = np.array(num2)
num3 = np.array(num3)
num4 = np.array(num4)

```

```

In [19]: # reshapping nparray
num1 = num1.reshape((num1.shape[0], 1))
num2 = num2.reshape((num2.shape[0], 1))
num3 = num3.reshape((num3.shape[0], 1))
num4 = num4.reshape((num4.shape[0], 1))

```

```

In [20]: # reshapping nparray
rev1 = rev1.reshape((rev1.shape[0], 1))
rev2 = rev2.reshape((rev2.shape[0], 1))
rev3 = rev3.reshape((rev3.shape[0], 1))
rev4 = rev4.reshape((rev4.shape[0], 1))

```

```

In [21]: # adding tfidf scores to dataframes
df1 = pd.DataFrame(np.hstack((num1, rev1)), columns = ['tfidf', 'review_body'])
df2 = pd.DataFrame(np.hstack((num2, rev2)), columns = ['tfidf', 'review_body'])
df3 = pd.DataFrame(np.hstack((num3, rev3)), columns = ['tfidf', 'review_body'])
df4 = pd.DataFrame(np.hstack((num4, rev4)), columns = ['tfidf', 'review_body'])

```

```
In [22]: # converts tfidf vectors to numeric form
df1['tfidf'] = pd.to_numeric(df1['tfidf'])
df2['tfidf'] = pd.to_numeric(df2['tfidf'])
df3['tfidf'] = pd.to_numeric(df3['tfidf'])
df4['tfidf'] = pd.to_numeric(df4['tfidf'])

In [23]: # getting length of reviews
df1['Length'] = df1['review_body'].str.len()
df2['Length'] = df2['review_body'].str.len()
df3['Length'] = df3['review_body'].str.len()
df4['Length'] = df4['review_body'].str.len()

In [24]: # ranking reviews according to how many times the words have appeared in a sent
df1['rank'] = df1['tfidf']/df1['Length']
df2['rank'] = df2['tfidf']/df2['Length']
df3['rank'] = df3['tfidf']/df3['Length']
df4['rank'] = df4['tfidf']/df4['Length']

In [25]: # sorting them according to the ranks
df1 = df1.sort_values(by='rank', ascending=False)
df2 = df2.sort_values(by='rank', ascending=False)
df3 = df3.sort_values(by='rank', ascending=False)
df4 = df4.sort_values(by='rank', ascending=False)

In [26]: # sub sampling the highest 20000 ranks from each rating
df1 = df1.head(20000)
df2 = df2.head(20000)
df3 = df3.head(20000)
df4 = df4.head(20000)

In [27]: # removing all unrequired columns for models in the later part of the project
df1 = pd.DataFrame(df1['review_body']).reset_index(drop=True)
df2 = pd.DataFrame(df2['review_body']).reset_index(drop=True)
df3 = pd.DataFrame(df3['review_body']).reset_index(drop=True)
df4 = pd.DataFrame(df4['review_body']).reset_index(drop=True)

In [28]: # creating labels lists

labels1 = [1]*20000
labels2 = [2]*20000
labels3 = [3]*20000
labels4 = [4]*20000
labels5 = [5]*20000

In [29]: # converting from list to DF
l1,l2,l3,l4,l5 = pd.DataFrame(), pd.DataFrame(), pd.DataFrame(), pd.DataFrame()
l1['star_rating'] = pd.DataFrame(labels1)
l2['star_rating'] = pd.DataFrame(labels2)
l3['star_rating'] = pd.DataFrame(labels3)
l4['star_rating'] = pd.DataFrame(labels4)
l5['star_rating'] = pd.DataFrame(labels5)

In [30]: df5 = pd.DataFrame(rating5['review_body']).reset_index(drop=True)

In [31]: # combining all dataframes
frame_combined = pd.concat([df1,df2,df3,df4,df5])
```

```
label_combined = pd.concat([l1,l2,l3,l4,l5])
combined = pd.concat([frame_combined, label_combined], axis = 1)
```

```
In [32]: # data = pd.concat([df['star_rating'], df['review_body']], axis=1)
# data = data.dropna()
# #data.drop(data[data['review_body'].str.split().str.len() < 10].index, inplace=True)
# data['star_rating'] = [int(i) for i in data['star_rating']]
# rating1 = data[data['star_rating'] == 1].sample(n=20000, random_state = 2)
# rating2 = data[data['star_rating'] == 2].sample(n=20000, random_state = 2)
# rating3 = data[data['star_rating'] == 3].sample(n=20000, random_state = 2)
# rating4 = data[data['star_rating'] == 4].sample(n=20000, random_state = 2)
# rating5 = data[data['star_rating'] == 5].sample(n=20000, random_state = 2)
# combined = pd.concat([rating1, rating2, rating3, rating4, rating5])
# combined
```

2. Word Embedding

(a) Load the pretrained "word2vec-google-news-300" Word2Vec model and learn how to extract word embeddings for your dataset. Try to check semantic similarities of the generated vectors using three examples of your own, e.g., King – Man + Woman = Queen or excellent ~ outstanding.

```
In [33]: # downloading all google word2vec vectors
import gensim.downloader as api
googlew2v = api.load('word2vec-google-news-300')
```

```
In [34]: #checking similarity between the 2 words using built in function
print("Most similar word and its similarity value: ", googlew2v.most_similar('golden_retriever'))

Most similar word and its similarity value: [('golden_retriever', 0.8104889392852783)]
```

```
In [35]: print("Similarity: ", googlew2v.similarity('fantastic', 'amazing'))

Similarity: 0.77898705
```

```
In [36]: print("Similarity: ", googlew2v.similarity('king', 'prince'))

Similarity: 0.61599934
```

```
In [37]: print("Top 3 similar words and their similarity values: ", googlew2v.most_similar('camry'))

Top 3 similar words and their similarity values: [('camry', 0.6310814023017883), ('chevy', 0.6251167058944702), ('camaro', 0.6154221892356873)]
```

```
In [38]: # example: Family - boy + girl = mother
# calculating cosine similarity and then using it to calculate similarity between
subtraction = googlew2v['family'] - googlew2v['boy']
addition = subtraction + googlew2v['girl']
num = np.dot(addition, googlew2v['mother'])
denom = (np.linalg.norm(addition) * np.linalg.norm(googlew2v['mother']))
dist_google = num/denom
dist_google
```

```
Out[38]: 0.5986675
```


(b) Train a Word2Vec model using your own dataset. You will use these extracted features in the subsequent questions of this assignment. Set the embedding size to be 300 and the window size to be 11. You can also consider a minimum word count of 10. Check the semantic similarities for the same two examples in part (a). What do you conclude from comparing vectors generated by yourself and the pretrained model? Which of the Word2Vec models seems to encode semantic similarities between words better? For the rest of this assignment, use the pretrained "word2vec-googlenews-300" Word2Vec features.

```
In [39]: # taken from reference given in hw2 pdf
from gensim import utils
from gensim.test.utils import datapath
combined['review_body'] = [str(i) for i in combined['review_body']]
class MyCorpus:
    def __iter__(self):
        corpus_path = datapath('lee_background.cor')
        for line in combined['review_body']:
            yield utils.simple_preprocess(line)

In [40]: # training own model
from gensim.models import Word2Vec
w2v = Word2Vec(sentences=MyCorpus(), vector_size=300, window=11, min_count=10)

In [41]: #checking similarity between the 2 words using built in function
print("Most similar word and its similarity value (own Word2Vec model): ", w2v.wv.most_similar('woman'))

Most similar word and its similarity value (own Word2Vec model): [('woman', 0.7054744362831116)]

In [42]: print("Most similar word and its similarity value (Google news): ", googlew2v.wv.most_similar('woman'))

Most similar word and its similarity value (Google news): [('queen', 0.7118193507194519)]

In [43]: print("Similarity value (own Word2Vec model): ", w2v.wv.similarity('sun', 'moon'))

Similarity value (own Word2Vec model): 0.49703205

In [44]: print("Similarity value (Google news): ", googlew2v.wv.similarity('sun', 'moon'))

Similarity value (Google news): 0.42628342

In [45]: # example: Family - boy + girl = mother
# calculating cosine similarity and then using it to calculate similarity between
subtraction = w2v.wv['family'] - w2v.wv['boy']
addition = subtraction + w2v.wv['girl']
num2 = np.dot(addition, w2v.wv['mother'])
denom2 = (np.linalg.norm(addition) * np.linalg.norm(w2v.wv['mother']))
dist_own = num2/denom2
dist_own

Out[45]: 0.6756631
```

The similarity values for my own model are sometimes less or more than the pretrained model depending on the example but there are a lot of vocabulary missing from my own model. Moreover, the vectors of words in my own model sometimes makes no sense. The pretrained model gives better "most similar" words.

3. Simple models

Using the Google pre-trained Word2Vec features, train a perceptron and an SVM model for the five class classification problem. For this purpose, use the average Word2Vec vectors for each review as the input feature ($x = \frac{1}{N} \sum_{i=1}^N \text{word}_i$ for review with N words). Report your accuracy values on the testing split for these models similar to HW1, i.e., for each of perceptron and SVM models, report two accuracy values Word2Vec and TF-IDF features. What do you conclude from comparing performances for the models trained using the two different feature types (TF-IDF and your trained

TFIDF

```
In [46]: # converted all ratings to int
combined['star_rating'] = [int(i) for i in combined['star_rating']]

In [47]: # using TFIDF vectorization for feature extraction

vectorizer = TfidfVectorizer()
fitt = vectorizer.fit_transform(combined['review_body'])

In [48]: # splitting the data as 80% training data and 20% testing data
xtrain_tfidf, xtest_tfidf, ytrain_tfidf, ytest_tfidf = train_test_split(fitt, combined['star_rating'],
                                test_size=0.2, random_state=0)

In [49]: # training pretrained TFIDF vectors through Perceptron model
from sklearn.linear_model import Perceptron
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
tfidf_perc_mod = Perceptron(random_state=0)

In [50]: tfidf_perc_mod.fit(xtrain_tfidf, ytrain_tfidf)
pred_perc_tfidf = tfidf_perc_mod.predict(xtest_tfidf)

In [51]: print('Accuracy using the TFIDF(Perceptron):', accuracy_score(ytest_tfidf, pred_perc_tfidf))
Accuracy using the TFIDF(Perceptron): 0.569

In [52]: # training pretrained TFIDF vectors through LinearSVC model
from sklearn.svm import LinearSVC
tfidf_svm_mod = LinearSVC(max_iter=1000, random_state=0)

In [53]: tfidf_svm_mod.fit(xtrain_tfidf, ytrain_tfidf)
pred_svm_tfidf = tfidf_svm_mod.predict(xtest_tfidf)

In [54]: print('Accuracy using the TFIDF(SVM):', accuracy_score(ytest_tfidf, pred_svm_tfidf))
Accuracy using the TFIDF(SVM): 0.65085
```

Word2Vec

```
In [55]: # function to average Word2Vec vectors for each review and tackling NaN results
def average(x, w2v):
    count, summ, ty = 0, np.zeros(shape=(300,)), type(x)
    if ty == str:
        lst = x.split(' ')
    elif ty == list:
        lst = x
    for w in lst:
        if w in w2v:
            word = w2v[w]
            count, summ = count + 1, summ + word
    if count == 0:
        return summ
    else:
        return (summ / count)
```

```
In [56]: #preparing Word2Vec vector for splitting
w2vxdata = combined['review_body'].apply(lambda x: average(x, googlew2v))
w2vxdata = np.array(w2vxdata.values.tolist())
w2vydata = combined['star_rating']
w2vydata = np.array(w2vydata.values.tolist())
```

```
In [57]: # splitting the data as 80% training data and 20% testing data
xtrain_google, xtest_google, ytrain_google, ytest_google = train_test_split(w2vxdata, w2vydata, test_size=0.2, random_state=0)
```

```
In [58]: # training pretrained Word2Vec embedding through Perceptron model
google_perc_mod = Perceptron(random_state=0)
google_perc_mod.fit(xtrain_google, ytrain_google)
```

```
Out[58]: Perceptron()
```

```
In [59]: pred_perc_google = google_perc_mod.predict(xtest_google)
```

```
In [60]: print('Accuracy using the pretrained Word2Vec model(Perceptron):', accuracy_score(ytest_google, pred_perc_google))
Accuracy using the pretrained Word2Vec model(Perceptron): 0.4835
```

```
In [61]: # training pretrained Word2Vec embedding through Perceptron model
google_svm_mod = LinearSVC(max_iter=1000, random_state=0)
```

```
In [62]: google_svm_mod.fit(xtrain_google, ytrain_google)
pred_svm_google = google_svm_mod.predict(xtest_google)
```

```
In [63]: print('Accuracy using the pretrained Word2Vec model(SVM):', accuracy_score(ytest_google, pred_svm_google))
Accuracy using the pretrained Word2Vec model(SVM): 0.5945
Accuracy for the TFIDF vectorization is more than Word2Vec model
```

4. Feedforward Neural Networks

Using the Word2Vec features, train a feedforward multilayer perceptron network for classification. Consider a network with two hidden layers, each with 50 and 10 nodes, respectively. You can use cross entropy loss and your own choice for other hyperparameters,

e.g., nonlinearity, number of epochs, etc. Part of getting good results is to select good values for these hyperparameters. You can also refer to the following tutorial to familiarize yourself: <https://www.kaggle.com/mishra1993/pytorch-multi-layer-perceptron-mnist> Although the above tutorial is for image data but the concept of training an MLP is very similar to what we want to do.

(a) To generate the input features, use the average Word2Vec vectors similar to the "Simple models" section and train the neural network. Report accuracy values on the testing split for your MLP.

```
In [64]: import torch
from torch.utils.data import DataLoader, Dataset
from torch.utils.data.sampler import SubsetRandomSampler
import matplotlib.pyplot as plt
```

```
In [65]: # custom dataset
# we have to overwrite len() and getitem() functions
class TrainDataset(Dataset):

    def __init__(self, xtrain, ytrain):
        self.data = xtrain
        self.labels = ytrain

    def __len__(self):
        return len(self.data)

    def __getitem__(self, index):
        data = self.data[index]
        label = self.labels[index]
        return data, label
```

```
In [66]: class TestDataset(Dataset):

    def __init__(self, xtest, ytest):
        self.data = xtest
        self.labels = ytest

    def __len__(self):
        return len(self.data)

    def __getitem__(self, index):
        data = self.data[index]
        label = self.labels[index]
        return data, label
```

```
In [67]: train_data = TrainDataset(xtrain_google, ytrain_google-1)
test_data = TestDataset(xtest_google, ytest_google-1)
```

```
In [68]: # initialising batch_size, valid_size
num_workers, batch_size, valid_size = 0, 32, 0.2
# converting data to torch.FloatTensor
# obtain training indices that will be used for validation
num_train = len(train_data)
indices = list(range(num_train))
np.random.shuffle(indices)
```

```

split = int(np.floor(valid_size * num_train))
train_idx, valid_idx = indices[split:], indices[:split]

# defining samplers for obtaining training and validation batches
train_sampler, valid_sampler = SubsetRandomSampler(train_idx), SubsetRandomSampler(valid_idx)

# prepare data loaders
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size, shuffle=True)
valid_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size, shuffle=False)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size, shuffle=False)

```

```

In [69]: # custom FNN
import torch.nn as nn
import torch.nn.functional as F

# define the NN architecture
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        # number of hidden nodes in each layer (512)
        hidden_1 = 50
        hidden_2 = 10
        # linear layer 1
        self.fc1 = nn.Linear(300, hidden_1)
        # linear layer 2
        self.fc2 = nn.Linear(hidden_1, hidden_2)
        # linear layer 3
        self.fc3 = nn.Linear(hidden_2, 5)
        # dropout layer (p=0.2)
        # dropout prevents overfitting of data
        self.dropout = nn.Dropout(0.1)

    def forward(self, x):
        # flatten image input
        x = x.view(-1, 300)
        # add hidden layer, with relu activation function
        x = F.relu(self.fc1(x))
        # add dropout layer
        x = self.dropout(x)
        # add hidden layer, with relu activation function
        x = F.relu(self.fc2(x))
        # add dropout layer
        x = self.dropout(x)
        # add output layer
        x = self.fc3(x)
        return x

# initialize the NN
model = Net()

```

```

In [70]: # loss function used
criterion = nn.CrossEntropyLoss()
# optimizer used
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)

```

```

In [71]: n_epochs, valid_loss_min = 100, np.Inf

for epoch in range(n_epochs):
    # monitor training loss

```

```

train_loss, valid_loss = 0.0, 0.0

# training model
model.train() # prep model for training
for data, target in train_loader:
    # clear the gradients of all optimized variables
    optimizer.zero_grad()
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data.float())
    # calculate the loss
    loss = criterion(output, target)
    # backward pass: compute gradient of the loss with respect to model parameters
    loss.backward()
    # perform a single optimization step (parameter update)
    optimizer.step()
    # update running training loss
    train_loss = train_loss + loss.item()*data.size(0)

# validating model
model.eval() # prep model for evaluation
for data, target in valid_loader:
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data.float())
    # calculate the loss
    loss = criterion(output, target)
    # update running validation loss
    valid_loss = valid_loss + loss.item()*data.size(0)

# print training/validation statistics
# calculate average loss over an epoch
train_loss, valid_loss = train_loss/len(train_loader.dataset), valid_loss/len(valid_loader.dataset)
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(epoch, train_loss, valid_loss))

# save model if validation loss has decreased
if valid_loss <= valid_loss_min:
    print('Validation loss decreased {:.6f} --> {:.6f}). Saving model ...')
    torch.save(model.state_dict(), 'model.pt')
    valid_loss_min = valid_loss

```

```
Epoch: 1      Training Loss: 1.283530      Validation Loss: 0.317319
Validation loss decreased (inf --> 0.317319). Saving model ...
Epoch: 2      Training Loss: 1.207887      Validation Loss: 0.283512
Validation loss decreased (0.317319 --> 0.283512). Saving model ...
Epoch: 3      Training Loss: 1.098549      Validation Loss: 0.259687
Validation loss decreased (0.283512 --> 0.259687). Saving model ...
Epoch: 4      Training Loss: 1.031042      Validation Loss: 0.248246
Validation loss decreased (0.259687 --> 0.248246). Saving model ...
Epoch: 5      Training Loss: 1.000458      Validation Loss: 0.242145
Validation loss decreased (0.248246 --> 0.242145). Saving model ...
Epoch: 6      Training Loss: 0.972072      Validation Loss: 0.231757
Validation loss decreased (0.242145 --> 0.231757). Saving model ...
Epoch: 7      Training Loss: 0.932853      Validation Loss: 0.222425
Validation loss decreased (0.231757 --> 0.222425). Saving model ...
Epoch: 8      Training Loss: 0.908597      Validation Loss: 0.217585
Validation loss decreased (0.222425 --> 0.217585). Saving model ...
Epoch: 9      Training Loss: 0.891317      Validation Loss: 0.213544
Validation loss decreased (0.217585 --> 0.213544). Saving model ...
Epoch: 10     Training Loss: 0.877228      Validation Loss: 0.209856
Validation loss decreased (0.213544 --> 0.209856). Saving model ...
Epoch: 11     Training Loss: 0.862411      Validation Loss: 0.206636
Validation loss decreased (0.209856 --> 0.206636). Saving model ...
Epoch: 12     Training Loss: 0.852515      Validation Loss: 0.203664
Validation loss decreased (0.206636 --> 0.203664). Saving model ...
Epoch: 13     Training Loss: 0.842260      Validation Loss: 0.201181
Validation loss decreased (0.203664 --> 0.201181). Saving model ...
Epoch: 14     Training Loss: 0.830667      Validation Loss: 0.198034
Validation loss decreased (0.201181 --> 0.198034). Saving model ...
Epoch: 15     Training Loss: 0.821941      Validation Loss: 0.195727
Validation loss decreased (0.198034 --> 0.195727). Saving model ...
Epoch: 16     Training Loss: 0.811856      Validation Loss: 0.193285
Validation loss decreased (0.195727 --> 0.193285). Saving model ...
Epoch: 17     Training Loss: 0.804343      Validation Loss: 0.191529
Validation loss decreased (0.193285 --> 0.191529). Saving model ...
Epoch: 18     Training Loss: 0.798134      Validation Loss: 0.189489
Validation loss decreased (0.191529 --> 0.189489). Saving model ...
Epoch: 19     Training Loss: 0.789906      Validation Loss: 0.187667
Validation loss decreased (0.189489 --> 0.187667). Saving model ...
Epoch: 20     Training Loss: 0.786103      Validation Loss: 0.186551
Validation loss decreased (0.187667 --> 0.186551). Saving model ...
Epoch: 21     Training Loss: 0.778144      Validation Loss: 0.185078
Validation loss decreased (0.186551 --> 0.185078). Saving model ...
Epoch: 22     Training Loss: 0.774564      Validation Loss: 0.183956
Validation loss decreased (0.185078 --> 0.183956). Saving model ...
Epoch: 23     Training Loss: 0.768200      Validation Loss: 0.182965
Validation loss decreased (0.183956 --> 0.182965). Saving model ...
Epoch: 24     Training Loss: 0.764111      Validation Loss: 0.182040
Validation loss decreased (0.182965 --> 0.182040). Saving model ...
Epoch: 25     Training Loss: 0.761840      Validation Loss: 0.181428
Validation loss decreased (0.182040 --> 0.181428). Saving model ...
Epoch: 26     Training Loss: 0.757878      Validation Loss: 0.180491
Validation loss decreased (0.181428 --> 0.180491). Saving model ...
Epoch: 27     Training Loss: 0.753142      Validation Loss: 0.179846
Validation loss decreased (0.180491 --> 0.179846). Saving model ...
Epoch: 28     Training Loss: 0.749955      Validation Loss: 0.178730
Validation loss decreased (0.179846 --> 0.178730). Saving model ...
Epoch: 29     Training Loss: 0.747823      Validation Loss: 0.178299
Validation loss decreased (0.178730 --> 0.178299). Saving model ...
Epoch: 30     Training Loss: 0.744259      Validation Loss: 0.177401
Validation loss decreased (0.178299 --> 0.177401). Saving model ...
```

```
Epoch: 31      Training Loss: 0.741140      Validation Loss: 0.176680
Validation loss decreased (0.177401 --> 0.176680). Saving model ...
Epoch: 32      Training Loss: 0.739769      Validation Loss: 0.176523
Validation loss decreased (0.176680 --> 0.176523). Saving model ...
Epoch: 33      Training Loss: 0.736456      Validation Loss: 0.175965
Validation loss decreased (0.176523 --> 0.175965). Saving model ...
Epoch: 34      Training Loss: 0.735313      Validation Loss: 0.175633
Validation loss decreased (0.175965 --> 0.175633). Saving model ...
Epoch: 35      Training Loss: 0.732908      Validation Loss: 0.174756
Validation loss decreased (0.175633 --> 0.174756). Saving model ...
Epoch: 36      Training Loss: 0.728682      Validation Loss: 0.174300
Validation loss decreased (0.174756 --> 0.174300). Saving model ...
Epoch: 37      Training Loss: 0.727095      Validation Loss: 0.174149
Validation loss decreased (0.174300 --> 0.174149). Saving model ...
Epoch: 38      Training Loss: 0.725503      Validation Loss: 0.173683
Validation loss decreased (0.174149 --> 0.173683). Saving model ...
Epoch: 39      Training Loss: 0.724420      Validation Loss: 0.173942
Epoch: 40      Training Loss: 0.721939      Validation Loss: 0.173747
Epoch: 41      Training Loss: 0.720678      Validation Loss: 0.172963
Validation loss decreased (0.173683 --> 0.172963). Saving model ...
Epoch: 42      Training Loss: 0.719394      Validation Loss: 0.172775
Validation loss decreased (0.172963 --> 0.172775). Saving model ...
Epoch: 43      Training Loss: 0.717943      Validation Loss: 0.173257
Epoch: 44      Training Loss: 0.716050      Validation Loss: 0.172351
Validation loss decreased (0.172775 --> 0.172351). Saving model ...
Epoch: 45      Training Loss: 0.714429      Validation Loss: 0.172020
Validation loss decreased (0.172351 --> 0.172020). Saving model ...
Epoch: 46      Training Loss: 0.712585      Validation Loss: 0.172018
Validation loss decreased (0.172020 --> 0.172018). Saving model ...
Epoch: 47      Training Loss: 0.710972      Validation Loss: 0.171357
Validation loss decreased (0.172018 --> 0.171357). Saving model ...
Epoch: 48      Training Loss: 0.709417      Validation Loss: 0.171162
Validation loss decreased (0.171357 --> 0.171162). Saving model ...
Epoch: 49      Training Loss: 0.708836      Validation Loss: 0.171316
Epoch: 50      Training Loss: 0.706041      Validation Loss: 0.171605
Epoch: 51      Training Loss: 0.704654      Validation Loss: 0.170363
Validation loss decreased (0.171162 --> 0.170363). Saving model ...
Epoch: 52      Training Loss: 0.704432      Validation Loss: 0.170718
Epoch: 53      Training Loss: 0.704036      Validation Loss: 0.170554
Epoch: 54      Training Loss: 0.700363      Validation Loss: 0.170224
Validation loss decreased (0.170363 --> 0.170224). Saving model ...
Epoch: 55      Training Loss: 0.700776      Validation Loss: 0.170178
Validation loss decreased (0.170224 --> 0.170178). Saving model ...
Epoch: 56      Training Loss: 0.699606      Validation Loss: 0.170063
Validation loss decreased (0.170178 --> 0.170063). Saving model ...
Epoch: 57      Training Loss: 0.699103      Validation Loss: 0.169592
Validation loss decreased (0.170063 --> 0.169592). Saving model ...
Epoch: 58      Training Loss: 0.698719      Validation Loss: 0.170323
Epoch: 59      Training Loss: 0.697625      Validation Loss: 0.169342
Validation loss decreased (0.169592 --> 0.169342). Saving model ...
Epoch: 60      Training Loss: 0.695581      Validation Loss: 0.169525
Epoch: 61      Training Loss: 0.692989      Validation Loss: 0.168990
Validation loss decreased (0.169342 --> 0.168990). Saving model ...
Epoch: 62      Training Loss: 0.694320      Validation Loss: 0.168724
Validation loss decreased (0.168990 --> 0.168724). Saving model ...
Epoch: 63      Training Loss: 0.692491      Validation Loss: 0.168858
Epoch: 64      Training Loss: 0.692057      Validation Loss: 0.168538
Validation loss decreased (0.168724 --> 0.168538). Saving model ...
Epoch: 65      Training Loss: 0.691134      Validation Loss: 0.168617
Epoch: 66      Training Loss: 0.689283      Validation Loss: 0.168328
```



```

Validation loss decreased (0.168538 --> 0.168328). Saving model ...
Epoch: 67      Training Loss: 0.686578      Validation Loss: 0.169276
Epoch: 68      Training Loss: 0.689837      Validation Loss: 0.167913
Validation loss decreased (0.168328 --> 0.167913). Saving model ...
Epoch: 69      Training Loss: 0.687044      Validation Loss: 0.168103
Epoch: 70      Training Loss: 0.686801      Validation Loss: 0.168339
Epoch: 71      Training Loss: 0.686725      Validation Loss: 0.168357
Epoch: 72      Training Loss: 0.684611      Validation Loss: 0.168353
Epoch: 73      Training Loss: 0.684309      Validation Loss: 0.168199
Epoch: 74      Training Loss: 0.683142      Validation Loss: 0.167983
Epoch: 75      Training Loss: 0.680762      Validation Loss: 0.168337
Epoch: 76      Training Loss: 0.680892      Validation Loss: 0.167249
Validation loss decreased (0.167913 --> 0.167249). Saving model ...
Epoch: 77      Training Loss: 0.680362      Validation Loss: 0.167816
Epoch: 78      Training Loss: 0.680028      Validation Loss: 0.167316
Epoch: 79      Training Loss: 0.679358      Validation Loss: 0.167248
Validation loss decreased (0.167249 --> 0.167248). Saving model ...
Epoch: 80      Training Loss: 0.678520      Validation Loss: 0.167224
Validation loss decreased (0.167248 --> 0.167224). Saving model ...
Epoch: 81      Training Loss: 0.677902      Validation Loss: 0.167936
Epoch: 82      Training Loss: 0.677249      Validation Loss: 0.167007
Validation loss decreased (0.167224 --> 0.167007). Saving model ...
Epoch: 83      Training Loss: 0.675521      Validation Loss: 0.168096
Epoch: 84      Training Loss: 0.675564      Validation Loss: 0.167307
Epoch: 85      Training Loss: 0.674490      Validation Loss: 0.166934
Validation loss decreased (0.167007 --> 0.166934). Saving model ...
Epoch: 86      Training Loss: 0.672709      Validation Loss: 0.167865
Epoch: 87      Training Loss: 0.673517      Validation Loss: 0.166742
Validation loss decreased (0.166934 --> 0.166742). Saving model ...
Epoch: 88      Training Loss: 0.674649      Validation Loss: 0.167214
Epoch: 89      Training Loss: 0.673019      Validation Loss: 0.168255
Epoch: 90      Training Loss: 0.671031      Validation Loss: 0.166525
Validation loss decreased (0.166742 --> 0.166525). Saving model ...
Epoch: 91      Training Loss: 0.670820      Validation Loss: 0.167938
Epoch: 92      Training Loss: 0.670605      Validation Loss: 0.166804
Epoch: 93      Training Loss: 0.668151      Validation Loss: 0.166818
Epoch: 94      Training Loss: 0.668928      Validation Loss: 0.167139
Epoch: 95      Training Loss: 0.668372      Validation Loss: 0.168603
Epoch: 96      Training Loss: 0.666899      Validation Loss: 0.167352
Epoch: 97      Training Loss: 0.666013      Validation Loss: 0.166048
Validation loss decreased (0.166525 --> 0.166048). Saving model ...
Epoch: 98      Training Loss: 0.665806      Validation Loss: 0.169283
Epoch: 99      Training Loss: 0.666535      Validation Loss: 0.166890
Epoch: 100     Training Loss: 0.664908      Validation Loss: 0.166193

```

```

In [72]: # loading the model with lowest validation loss
         model.load_state_dict(torch.load('model.pt'))

```

```

Out[72]: <All keys matched successfully>

```

```

In [73]: correct, total = 0, 0
         # since we're not training, we don't need to calculate the gradients for out ou
         with torch.no_grad():
             for data in test_loader:
                 embeddings, labels = data
                 # calculating outputs by running embeddings through the network
                 model.to("cpu")
                 outputs = model(embeddings.float())
                 # the class with the highest score is what we choose as prediction

```

```
_, predicted = torch.max(outputs.data, 1)
total = total + labels.size(0)
correct = correct + (predicted == labels).sum().item()
```

```
In [74]: print('Accuracy for Average Word2Vec vectors FNN model: ', (100 * correct / tot

Accuracy for Average Word2Vec vectors FNN model: 66.3
```

(b) To generate the input features, concatenate the first 10 Word2Vec vectors for each review as the input feature ($x = [WT\ 1, \dots, WT\ 10]$) and train the neural network. Report the accuracy value on the testing split for your MLP model. What do you conclude by comparing accuracy values you obtain with those obtained in the "Simple Models" section.

```
In [85]: # Function to generate the input features, concatenate the first 10 Word2Vec ve
def concatenation(s, w2v):
    a, b = 0, 0
    if type(s) == list:
        lst = s
    if type(s) == str:
        lst = s.split(' ')
    leng = len(lst)
    while (a < leng) & (b < 10):
        if lst[a] in w2v:
            wv = w2v[lst[a]]
            if b != 0:
                res = np.concatenate((res, wv))
            else:
                res = wv
            a, b = a + 1, b + 1
            #print("a: ", a)
            #print("b: ", b)
        else:
            a = a + 1
    #print("b final:", b)
    if b < 10:
        if (10 - b) != 10:
            res = np.concatenate((res, np.zeros(shape=(300*(10 - b), ))))
        else:
            res = np.zeros(shape=(300*(10 - b), ))
    return res
```

```
In [86]: # preparing the Word2Vec vectors for splitting
w2vx10data = combined['review_body'].apply(lambda x: concatenation(x, googlew2v
w2vx10data = np.array(w2vx10data.values.tolist())
w2vy10data = combined['star_rating']
w2vy10data = np.array(w2vy10data.values.tolist())
```

```
In [87]: # splitting the data as 80% training data and 20% testing data
xtrain10_google, xtest10_google, ytrain10_google, ytest10_google = train_test_s
```

```
In [88]: # custom dataset
class TrainDataset(Dataset):

    def __init__(self, xtrain, ytrain):
        self.data = xtrain
        self.labels = ytrain
```

```

def __len__(self):
    return len(self.data)

def __getitem__(self, index):
    data = self.data[index]
    label = self.labels[index]
    return data, label

```

```

In [89]: class TestDataset(Dataset):

def __init__(self, xtest, ytest):
    self.data = xtest
    self.labels = ytest

def __len__(self):
    return len(self.data)

def __getitem__(self, index):
    data = self.data[index]
    label = self.labels[index]
    return data, label

```

```

In [90]: train10_data = TrainDataset(xtrain10_google, ytrain10_google-1)
test10_data = TestDataset(xtest10_google, ytest10_google-1)

```

```

In [91]: num_workers, batch_size, valid_size = 0, 32, 0.2

# obtain training indices that will be used for validation
num_train = len(train10_data)
indices = list(range(num_train))
np.random.shuffle(indices)
split = int(np.floor(valid_size * num_train))
train_idx, valid_idx = indices[split:], indices[:split]

# define samplers for obtaining training and validation batches
train_sampler, valid_sampler = SubsetRandomSampler(train_idx), SubsetRandomSampler(valid_idx)

# prepare data loaders
train_loader = torch.utils.data.DataLoader(train10_data, batch_size=batch_size,
valid_loader = torch.utils.data.DataLoader(train10_data, batch_size=batch_size,
test_loader = torch.utils.data.DataLoader(test10_data, batch_size=batch_size, r

```

```

In [92]: # custom FNN
import torch.nn as nn
import torch.nn.functional as F

# define the FNN architecture
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        # number of hidden nodes in each layer (512)
        hidden_1 = 50
        hidden_2 = 10
        # linear layer 1
        self.fc1 = nn.Linear(3000, hidden_1)
        # linear layer 2
        self.fc2 = nn.Linear(hidden_1, hidden_2)
        # linear layer 3

```

```

self.fc3 = nn.Linear(hidden_2, 5)
# dropout layer (p=0.2)
# dropout prevents overfitting of data
self.dropout = nn.Dropout(0.2)

def forward(self, x):
    # flatten image input
    x = x.view(-1, 3000)
    # add hidden layer, with relu activation function
    x = F.relu(self.fc1(x))
    # add dropout layer
    x = self.dropout(x)
    # add hidden layer, with relu activation function
    x = F.relu(self.fc2(x))
    # add dropout layer
    x = self.dropout(x)
    # add output layer
    x = self.fc3(x)
    return x

# initialize the NN
model10 = Net()

```

```

In [93]: # loss function used
criterion = nn.CrossEntropyLoss()
# optimizer used
optimizer = torch.optim.SGD(model10.parameters(), lr=0.0075)

```

```

In [94]: n_epochs, valid_loss_min = 30, np.Inf

for epoch in range(n_epochs):
    # monitor training loss
    train_loss, valid_loss = 0.0, 0.0
    # training model
    model10.train() # prep model for training
    for data, target in train_loader:
        # clear the gradients of all optimized variables
        optimizer.zero_grad()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model10(data.float())
        # calculate the loss
        loss = criterion(output, target)
        # backward pass: compute gradient of the loss with respect to model parameters
        loss.backward()
        # perform a single optimization step (parameter update)
        optimizer.step()
        # update running training loss
        train_loss = train_loss + loss.item()*data.size(0)

    # validating model
    model10.eval() # prep model for evaluation
    for data, target in valid_loader:
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model10(data.float())
        # calculate the loss
        loss = criterion(output, target)
        # update running validation loss
        valid_loss = valid_loss + loss.item()*data.size(0)

```

```
# print training/validation statistics
# calculate average loss over an epoch
train_loss, valid_loss = train_loss/len(train_loader.dataset), valid_loss/len(valid_loader.dataset)
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(epoch, train_loss, valid_loss))

# save model if validation loss has decreased
if valid_loss <= valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(valid_loss_min, valid_loss))
    torch.save(model10.state_dict(), 'model10.pt')
    valid_loss_min = valid_loss
```

```

Epoch: 1      Training Loss: 1.260506      Validation Loss: 0.296781
Validation loss decreased (inf --> 0.296781). Saving model ...
Epoch: 2      Training Loss: 1.044644      Validation Loss: 0.220601
Validation loss decreased (0.296781 --> 0.220601). Saving model ...
Epoch: 3      Training Loss: 0.896868      Validation Loss: 0.200710
Validation loss decreased (0.220601 --> 0.200710). Saving model ...
Epoch: 4      Training Loss: 0.845342      Validation Loss: 0.192679
Validation loss decreased (0.200710 --> 0.192679). Saving model ...
Epoch: 5      Training Loss: 0.813451      Validation Loss: 0.186653
Validation loss decreased (0.192679 --> 0.186653). Saving model ...
Epoch: 6      Training Loss: 0.790590      Validation Loss: 0.182170
Validation loss decreased (0.186653 --> 0.182170). Saving model ...
Epoch: 7      Training Loss: 0.775066      Validation Loss: 0.179425
Validation loss decreased (0.182170 --> 0.179425). Saving model ...
Epoch: 8      Training Loss: 0.761997      Validation Loss: 0.177023
Validation loss decreased (0.179425 --> 0.177023). Saving model ...
Epoch: 9      Training Loss: 0.750059      Validation Loss: 0.175439
Validation loss decreased (0.177023 --> 0.175439). Saving model ...
Epoch: 10     Training Loss: 0.739694      Validation Loss: 0.174131
Validation loss decreased (0.175439 --> 0.174131). Saving model ...
Epoch: 11     Training Loss: 0.731407      Validation Loss: 0.173073
Validation loss decreased (0.174131 --> 0.173073). Saving model ...
Epoch: 12     Training Loss: 0.722393      Validation Loss: 0.171998
Validation loss decreased (0.173073 --> 0.171998). Saving model ...
Epoch: 13     Training Loss: 0.716686      Validation Loss: 0.170993
Validation loss decreased (0.171998 --> 0.170993). Saving model ...
Epoch: 14     Training Loss: 0.708511      Validation Loss: 0.170373
Validation loss decreased (0.170993 --> 0.170373). Saving model ...
Epoch: 15     Training Loss: 0.701397      Validation Loss: 0.169133
Validation loss decreased (0.170373 --> 0.169133). Saving model ...
Epoch: 16     Training Loss: 0.696935      Validation Loss: 0.168920
Validation loss decreased (0.169133 --> 0.168920). Saving model ...
Epoch: 17     Training Loss: 0.691269      Validation Loss: 0.168499
Validation loss decreased (0.168920 --> 0.168499). Saving model ...
Epoch: 18     Training Loss: 0.684719      Validation Loss: 0.167780
Validation loss decreased (0.168499 --> 0.167780). Saving model ...
Epoch: 19     Training Loss: 0.679013      Validation Loss: 0.168175
Epoch: 20     Training Loss: 0.673193      Validation Loss: 0.167153
Validation loss decreased (0.167780 --> 0.167153). Saving model ...
Epoch: 21     Training Loss: 0.665874      Validation Loss: 0.166948
Validation loss decreased (0.167153 --> 0.166948). Saving model ...
Epoch: 22     Training Loss: 0.662770      Validation Loss: 0.166680
Validation loss decreased (0.166948 --> 0.166680). Saving model ...
Epoch: 23     Training Loss: 0.658224      Validation Loss: 0.166985
Epoch: 24     Training Loss: 0.653228      Validation Loss: 0.166784
Epoch: 25     Training Loss: 0.649802      Validation Loss: 0.166879
Epoch: 26     Training Loss: 0.646602      Validation Loss: 0.166523
Validation loss decreased (0.166680 --> 0.166523). Saving model ...
Epoch: 27     Training Loss: 0.640350      Validation Loss: 0.166935
Epoch: 28     Training Loss: 0.635417      Validation Loss: 0.166949
Epoch: 29     Training Loss: 0.631338      Validation Loss: 0.166895
Epoch: 30     Training Loss: 0.629316      Validation Loss: 0.166741

```

```

In [95]: # loading the model with lowest validation loss
         model10.load_state_dict(torch.load('model10.pt'))

```

```

Out[95]: <All keys matched successfully>

```

```

In [96]: correct, total = 0, 0

```

```
# since we're not training, we don't need to calculate the gradients for out ou
with torch.no_grad():
    for data in test_loader:
        embeddings, labels = data
        # calculating outputs by running embeddings through the network
        model10.to("cpu")
        outputs = model10(embeddings.float())
        # the class with the highest score is what we choose as prediction
        _, predicted = torch.max(outputs.data, 1)
        total = total + labels.size(0)
        correct = correct + (predicted == labels).sum().item()
```

```
In [97]: print('Accuracy for 10 Word2Vec vectors FNN model: ', (100 * correct / total))
```

Accuracy for 10 Word2Vec vectors FNN model: 67.015

5. Recurrent Neural Networks

Using the Word2Vec features, train a recurrent neural network (RNN) for classification. You can refer to the following tutorial to familiarize yourself:

https://pytorch.org/tutorials/intermediate/char_rnn_classification_tutorial.html (a) Train a simple RNN for sentiment analysis. You can consider an RNN cell with the hidden state size of 20. To feed your data into our RNN, limit the maximum review length to 20 by truncating longer reviews and padding shorter reviews with a null value (0). Report accuracy values on the testing split for your RNN model. What do you conclude by comparing accuracy values you obtain with those obtained with feedforward neural network models.

```
In [131... # function to feed your data into our RNN, limit the maximum review length to 20
#by truncating longer reviews and padding shorter reviews with a null value (0)
def review20(s, rnnw2v):
    a, b = 0, 0
    if type(s) == str:
        lst = s.split(' ')
    elif type(s) == list:
        lst = s
    leng = len(lst)
    while (b < 20) & (a < leng):
        if lst[a] in rnnw2v:
            if b != 0:
                w2v = rnnw2v[lst[a]]
                w2vm = np.vstack((w2vm, w2v))
            else:
                w2vm = rnnw2v[lst[a]]
            a, b = a + 1, b + 1
        else:
            a = a + 1
    if b != 0:
        zeros = np.zeros(shape=(20-b, 300))
        w2vm = np.vstack((w2vm, zeros))
    if b < 20:
        w2vm = np.zeros(shape=(20, 300))
    return w2vm
```

```
In [132... w2vxdatarnn = combined2['review_body'].apply(lambda x: review20(x, googlew2v))
```



```
w2vydatarnn = combined2['star_rating']
w2vxdatarnn = np.array(w2vxdatarnn.values.tolist())
w2vydatarnn = np.array(w2vydatarnn.values.tolist())
```

```
In [133... # splitting data into 80% training and 20% testing
xtrainrnn_google, xtestrnn_google, ytrainrnn_google, ytestrnn_google = train_te
```

```
In [134... # custom dataset
class TrainDataset(Dataset):

    def __init__(self, xtrain, ytrain):
        self.data = xtrain
        self.labels = ytrain

    def __len__(self):
        return len(self.data)

    def __getitem__(self, index):
        data = self.data[index]
        label = self.labels[index]

        return data, label
```

```
In [135... class TestDataset(Dataset):

    def __init__(self, xtest, ytest):
        self.data = xtest
        self.labels = ytest

    def __len__(self):
        return len(self.data)

    def __getitem__(self, index):
        data = self.data[index]
        label = self.labels[index]

        return data, label
```

```
In [136... train_datarnn = TrainDataset(xtrainrnn_google, ytrainrnn_google-1)
test_datarnn = TestDataset(xtestrnn_google, ytestrnn_google-1)
```

```
In [137... num_workers, batch_size, valid_size = 0, 32, 0.2
# convert data to torch.FloatTensor

# obtain training indices that will be used for validation
num_train = len(train_datarnn)
indices = list(range(num_train))
np.random.shuffle(indices)
split = int(np.floor(valid_size * num_train))
train_idx, valid_idx = indices[split:], indices[:split]

# define samplers for obtaining training and validation batches
train_sampler, valid_sampler = SubsetRandomSampler(train_idx), SubsetRandomSampler(valid_idx)

# prepare data loaders
train_loader_rnn = torch.utils.data.DataLoader(train_datarnn, batch_size=batch_size, shuffle=True)
valid_loader_rnn = torch.utils.data.DataLoader(train_datarnn, batch_size=batch_size, shuffle=False)
test_loader_rnn = torch.utils.data.DataLoader(test_datarnn, batch_size=batch_size, shuffle=False)
```

```

In [138... # custom RNN model
class RNNModel(nn.Module):
    def __init__(self, input_dim, hidden_dim, layer_dim, output_dim):
        super().__init__()

        # Number of hidden dimensions
        self.hidden_dim = hidden_dim
        # Number of hidden layers
        self.layer_dim = layer_dim
        # RNN
        self.rnn = nn.RNN(input_dim, hidden_dim, layer_dim, batch_first=True, r
        # Output layer
        self.fc = nn.Linear(hidden_dim, output_dim)

        self.softmax = nn.LogSoftmax(dim = 1)

    def forward(self, x):

        # Initialize hidden state with zeros
        h0 = torch.zeros(self.layer_dim, x.size(0), self.hidden_dim)
        # One time step
        out, hn = self.rnn(x, h0)
        out = self.fc(out[:, -1, :])
        out = self.softmax(out)
        return out

# initialize RNN
model_rnn = RNNModel(300, 20, 1, 5)

```

```

In [139... criterion = nn.NLLLoss()
optimizer_rnn = torch.optim.Adam(model_rnn.parameters(), lr=0.01)

```

```

In [140... # number of epochs to train the model
n_epochs, valid_loss_min = 30, np.Inf

for epoch in range(n_epochs):
    # monitor training loss
    train_loss, valid_loss = 0.0, 0.0

    # training the model
    model_rnn.train() # prep model for training
    for data, target in train_loader_rnn:
        # clear the gradients of all optimized variables
        optimizer_rnn.zero_grad()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model_rnn(data.float())
        # calculate the loss
        loss = criterion(output, target)
        # backward pass: compute gradient of the loss with respect to model parameters
        loss.backward()
        # perform a single optimization step (parameter update)
        optimizer_rnn.step()
        # update running training loss
        train_loss = train_loss + loss.item()*data.size(0)

    #validating model
    model_rnn.eval() # prep model for evaluation
    for data, target in valid_loader_rnn:

```

```

# forward pass: compute predicted outputs by passing inputs to the model
output = model_rnn(data.float())
# calculate the loss
loss = criterion(output, target)
# update running validation loss
valid_loss = valid_loss + loss.item()*data.size(0)

# print training/validation statistics
# calculate average loss over an epoch
train_loss, valid_loss = train_loss/len(train_loader_rnn.dataset), valid_loss/len(valid_loader_rnn.dataset)
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(epoch, train_loss, valid_loss))

# save model if validation loss has decreased
if valid_loss <= valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(valid_loss, valid_loss_min))
    torch.save(model_rnn.state_dict(), 'model.pt')
    valid_loss_min = valid_loss

```

```

Epoch: 1      Training Loss: 1.283001      Validation Loss: 0.317449
Validation loss decreased (inf --> 0.317449). Saving model ...
Epoch: 2      Training Loss: 1.252116      Validation Loss: 0.307635
Validation loss decreased (0.317449 --> 0.307635). Saving model ...
Epoch: 3      Training Loss: 1.217740      Validation Loss: 0.306188
Validation loss decreased (0.307635 --> 0.306188). Saving model ...
Epoch: 4      Training Loss: 1.199817      Validation Loss: 0.300398
Validation loss decreased (0.306188 --> 0.300398). Saving model ...
Epoch: 5      Training Loss: 1.198098      Validation Loss: 0.297399
Validation loss decreased (0.300398 --> 0.297399). Saving model ...
Epoch: 6      Training Loss: 1.189712      Validation Loss: 0.299164
Epoch: 7      Training Loss: 1.192799      Validation Loss: 0.297447
Epoch: 8      Training Loss: 1.189216      Validation Loss: 0.296797
Validation loss decreased (0.297399 --> 0.296797). Saving model ...
Epoch: 9      Training Loss: 1.184553      Validation Loss: 0.294477
Validation loss decreased (0.296797 --> 0.294477). Saving model ...
Epoch: 10     Training Loss: 1.194137      Validation Loss: 0.295632
Epoch: 11     Training Loss: 1.184758      Validation Loss: 0.298907
Epoch: 12     Training Loss: 1.180209      Validation Loss: 0.297092
Epoch: 13     Training Loss: 1.180336      Validation Loss: 0.295236
Epoch: 14     Training Loss: 1.177964      Validation Loss: 0.303544
Epoch: 15     Training Loss: 1.176563      Validation Loss: 0.296384
Epoch: 16     Training Loss: 1.177640      Validation Loss: 0.300024
Epoch: 17     Training Loss: 1.175896      Validation Loss: 0.299347
Epoch: 18     Training Loss: 1.177329      Validation Loss: 0.310685
Epoch: 19     Training Loss: 1.172469      Validation Loss: 0.295979
Epoch: 20     Training Loss: 1.176689      Validation Loss: 0.294404
Validation loss decreased (0.294477 --> 0.294404). Saving model ...
Epoch: 21     Training Loss: 1.175606      Validation Loss: 0.301625
Epoch: 22     Training Loss: 1.183474      Validation Loss: 0.314457
Epoch: 23     Training Loss: 1.197302      Validation Loss: 0.304831
Epoch: 24     Training Loss: 1.175817      Validation Loss: 0.299629
Epoch: 25     Training Loss: 1.172117      Validation Loss: 0.294882
Epoch: 26     Training Loss: 1.184060      Validation Loss: 0.298823
Epoch: 27     Training Loss: 1.171910      Validation Loss: 0.294755
Epoch: 28     Training Loss: 1.172655      Validation Loss: 0.298792
Epoch: 29     Training Loss: 1.170841      Validation Loss: 0.296271
Epoch: 30     Training Loss: 1.172918      Validation Loss: 0.297978

```

```

In [141]: # loading the model with lowest validation loss
model_rnn.load_state_dict(torch.load('model.pt'))

```

Out[141]: <All keys matched successfully>

```
In [142... correct, total = 0, 0
# since we're not training, we don't need to calculate the gradients for out ou
with torch.no_grad():
    for data in test_loader_rnn:
        embeddings, labels = data
        # calculating outputs by running embeddings through the network
        model_rnn.to("cpu")
        outputs = model_rnn(embeddings.float())
        # the class with the highest score is what we choose as prediction
        _, predicted = torch.max(outputs.data, 1)
        total = total + labels.size(0)
        correct = correct + (predicted == labels).sum().item()
```

```
In [143... print('Accuracy for RNN: ', (100 * correct / total))
```

Accuracy for RNN: 33.175

(b) Repeat part (a) by considering a gated recurrent unit cell. What do you conclude by comparing accuracy values you obtain with those obtained using simple RNN.

```
In [124... num_workers, batch_size, valid_size = 0, 32, 0.2

# convert data to torch.FloatTensor
# obtain training indices that will be used for validation
num_train = len(train_datarnn)
indices = list(range(num_train))
np.random.shuffle(indices)
split = int(np.floor(valid_size * num_train))
train_idx, valid_idx = indices[split:], indices[:split]

# define samplers for obtaining training and validation batches
train_sampler, valid_sampler = SubsetRandomSampler(train_idx), SubsetRandomSampler(valid_idx)
# prepare data loaders
train_loader_rnn = torch.utils.data.DataLoader(train_datarnn, batch_size=batch_size, num_workers=num_workers)
valid_loader_rnn = torch.utils.data.DataLoader(train_datarnn, batch_size=batch_size, num_workers=num_workers)
test_loader_rnn = torch.utils.data.DataLoader(test_datarnn, batch_size=batch_size, num_workers=num_workers)
```

```
In [125... # custom GatedRNN
class GatedRNN(nn.Module):
    def __init__(self, input_dim, hidden_dim, layer_dim, output_dim):
        super().__init__()

        # Number of hidden dimensions
        self.hidden_dim = hidden_dim
        # Number of hidden layers
        self.layer_dim = layer_dim
        # GRU
        self.gru = nn.GRU(input_dim, hidden_dim, layer_dim, batch_first=True)
        # Output layer
        self.fc = nn.Linear(hidden_dim, output_dim)
        self.softmax = nn.LogSoftmax(dim = 1)

    def forward(self, x):

        # Initialize hidden state with zeros
        h0 = torch.zeros(self.layer_dim, x.size(0), self.hidden_dim)
```

```

        # One time step
        out, hn = self.gru(x, h0)
        out = self.fc(out[:, -1, :])
        out = self.softmax(out)
        return out

model_gatedrnn = GatedRNN(300, 20, 1, 5)

```

```

In [126... # loss function used
criterion = nn.CrossEntropyLoss()
#optimizer used for gradient descent
optimizer_rnn = torch.optim.SGD(model_gatedrnn.parameters(), lr=0.0075)

```

```

In [127... # number of epochs to train the model
n_epochs, valid_loss_min = 50, np.Inf

for epoch in range(n_epochs):
    # monitor training loss
    train_loss, valid_loss = 0.0, 0.0
    #training model
    model_gatedrnn.train() # prep model for training
    for data, target in train_loader_rnn:
        # clear the gradients of all optimized variables
        optimizer_rnn.zero_grad()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model_gatedrnn(data.float())
        # calculate the loss
        loss = criterion(output, target)
        # backward pass: compute gradient of the loss with respect to model parameters
        loss.backward()
        # perform a single optimization step (parameter update)
        optimizer_rnn.step()
        # update running training loss
        train_loss = train_loss + loss.item()*data.size(0)

    # validating model
    model_gatedrnn.eval() # prep model for evaluation
    for data, target in valid_loader_rnn:
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model_gatedrnn(data.float())
        # calculate the loss
        loss = criterion(output, target)
        # update running validation loss
        valid_loss = valid_loss + loss.item()*data.size(0)

    # print training/validation statistics
    # calculate average loss over an epoch
    train_loss, valid_loss = train_loss/len(train_loader_rnn.dataset), valid_loss/len(valid_loader_rnn.dataset)
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(epoch, train_loss, valid_loss))

    # save model if validation loss has decreased
    if valid_loss <= valid_loss_min:
        print('Validation loss decreased {:.6f} --> {:.6f}). Saving model ...'.format(valid_loss, valid_loss_min))
        torch.save(model_gatedrnn.state_dict(), 'model.pt')
        valid_loss_min = valid_loss

```

```
Epoch: 1      Training Loss: 1.289548      Validation Loss: 0.321773
Validation loss decreased (inf --> 0.321773). Saving model ...
Epoch: 2      Training Loss: 1.285296      Validation Loss: 0.321020
Validation loss decreased (0.321773 --> 0.321020). Saving model ...
Epoch: 3      Training Loss: 1.281657      Validation Loss: 0.320157
Validation loss decreased (0.321020 --> 0.320157). Saving model ...
Epoch: 4      Training Loss: 1.277183      Validation Loss: 0.319084
Validation loss decreased (0.320157 --> 0.319084). Saving model ...
Epoch: 5      Training Loss: 1.270888      Validation Loss: 0.317087
Validation loss decreased (0.319084 --> 0.317087). Saving model ...
Epoch: 6      Training Loss: 1.257829      Validation Loss: 0.311675
Validation loss decreased (0.317087 --> 0.311675). Saving model ...
Epoch: 7      Training Loss: 1.232900      Validation Loss: 0.305766
Validation loss decreased (0.311675 --> 0.305766). Saving model ...
Epoch: 8      Training Loss: 1.217541      Validation Loss: 0.303540
Validation loss decreased (0.305766 --> 0.303540). Saving model ...
Epoch: 9      Training Loss: 1.207840      Validation Loss: 0.300587
Validation loss decreased (0.303540 --> 0.300587). Saving model ...
Epoch: 10     Training Loss: 1.199622      Validation Loss: 0.298878
Validation loss decreased (0.300587 --> 0.298878). Saving model ...
Epoch: 11     Training Loss: 1.192036      Validation Loss: 0.296773
Validation loss decreased (0.298878 --> 0.296773). Saving model ...
Epoch: 12     Training Loss: 1.184700      Validation Loss: 0.296729
Validation loss decreased (0.296773 --> 0.296729). Saving model ...
Epoch: 13     Training Loss: 1.176564      Validation Loss: 0.292547
Validation loss decreased (0.296729 --> 0.292547). Saving model ...
Epoch: 14     Training Loss: 1.168211      Validation Loss: 0.291987
Validation loss decreased (0.292547 --> 0.291987). Saving model ...
Epoch: 15     Training Loss: 1.161736      Validation Loss: 0.289605
Validation loss decreased (0.291987 --> 0.289605). Saving model ...
Epoch: 16     Training Loss: 1.156584      Validation Loss: 0.292754
Epoch: 17     Training Loss: 1.152186      Validation Loss: 0.287800
Validation loss decreased (0.289605 --> 0.287800). Saving model ...
Epoch: 18     Training Loss: 1.148037      Validation Loss: 0.289288
Epoch: 19     Training Loss: 1.145159      Validation Loss: 0.287960
Epoch: 20     Training Loss: 1.141931      Validation Loss: 0.285502
Validation loss decreased (0.287800 --> 0.285502). Saving model ...
Epoch: 21     Training Loss: 1.138719      Validation Loss: 0.284989
Validation loss decreased (0.285502 --> 0.284989). Saving model ...
Epoch: 22     Training Loss: 1.135665      Validation Loss: 0.289896
Epoch: 23     Training Loss: 1.133620      Validation Loss: 0.286784
Epoch: 24     Training Loss: 1.130854      Validation Loss: 0.283358
Validation loss decreased (0.284989 --> 0.283358). Saving model ...
Epoch: 25     Training Loss: 1.128849      Validation Loss: 0.283733
Epoch: 26     Training Loss: 1.125378      Validation Loss: 0.282474
Validation loss decreased (0.283358 --> 0.282474). Saving model ...
Epoch: 27     Training Loss: 1.123925      Validation Loss: 0.282077
Validation loss decreased (0.282474 --> 0.282077). Saving model ...
Epoch: 28     Training Loss: 1.121677      Validation Loss: 0.281045
Validation loss decreased (0.282077 --> 0.281045). Saving model ...
Epoch: 29     Training Loss: 1.119699      Validation Loss: 0.287675
Epoch: 30     Training Loss: 1.117810      Validation Loss: 0.281724
Epoch: 31     Training Loss: 1.115807      Validation Loss: 0.280709
Validation loss decreased (0.281045 --> 0.280709). Saving model ...
Epoch: 32     Training Loss: 1.113873      Validation Loss: 0.279994
Validation loss decreased (0.280709 --> 0.279994). Saving model ...
Epoch: 33     Training Loss: 1.112199      Validation Loss: 0.279635
Validation loss decreased (0.279994 --> 0.279635). Saving model ...
Epoch: 34     Training Loss: 1.110301      Validation Loss: 0.282371
Epoch: 35     Training Loss: 1.108387      Validation Loss: 0.281459
```

```

Epoch: 36      Training Loss: 1.106926      Validation Loss: 0.278647
Validation loss decreased (0.279635 --> 0.278647). Saving model ...
Epoch: 37      Training Loss: 1.105597      Validation Loss: 0.278015
Validation loss decreased (0.278647 --> 0.278015). Saving model ...
Epoch: 38      Training Loss: 1.103633      Validation Loss: 0.280401
Epoch: 39      Training Loss: 1.102408      Validation Loss: 0.281914
Epoch: 40      Training Loss: 1.100871      Validation Loss: 0.278361
Epoch: 41      Training Loss: 1.099510      Validation Loss: 0.276755
Validation loss decreased (0.278015 --> 0.276755). Saving model ...
Epoch: 42      Training Loss: 1.098190      Validation Loss: 0.277570
Epoch: 43      Training Loss: 1.096485      Validation Loss: 0.276254
Validation loss decreased (0.276755 --> 0.276254). Saving model ...
Epoch: 44      Training Loss: 1.095302      Validation Loss: 0.276883
Epoch: 45      Training Loss: 1.093840      Validation Loss: 0.275849
Validation loss decreased (0.276254 --> 0.275849). Saving model ...
Epoch: 46      Training Loss: 1.092995      Validation Loss: 0.275682
Validation loss decreased (0.275849 --> 0.275682). Saving model ...
Epoch: 47      Training Loss: 1.091815      Validation Loss: 0.275710
Epoch: 48      Training Loss: 1.090409      Validation Loss: 0.275230
Validation loss decreased (0.275682 --> 0.275230). Saving model ...
Epoch: 49      Training Loss: 1.089028      Validation Loss: 0.274874
Validation loss decreased (0.275230 --> 0.274874). Saving model ...
Epoch: 50      Training Loss: 1.088055      Validation Loss: 0.276718

```

```

In [128... # loading the model with the lowest validation loss
model_gatedrnn.load_state_dict(torch.load('model.pt'))

```

```

Out[128]: <All keys matched successfully>

```

```

In [129... correct, total = 0, 0
with torch.no_grad():
    for data in test_loader_rnn:
        embeddings, labels = data
        # calculating outputs by running embeddings through the network
        model_gatedrnn.to("cpu")
        outputs = model_gatedrnn(embeddings.float())
        # the class with the highest score is what we choose as prediction
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

```

```

In [130... print('Accuracy for Gated RNN model: ', (100 * correct / total))

```

Accuracy for Gated RNN model: 39.705

References:

https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

<https://www.kaggle.com/code/mishra1993/pytorch-multi-layer-perceptron-mnist/notebook>

https://pytorch.org/tutorials/intermediate/char_rnn_classification_tutorial.html

https://radimrehurek.com/gensim/auto_examples/tutorials/run_word2vec.html

In []:

In []:

